

From Controversy to Consensus: Modelling Community-Sensitive Common Ground Management in German Discourse Markers

This study investigates three German discourse markers, which we label *consensus markers*—*soviel ich weiß* ('as far as I know'), *ja* (roughly 'as you know'), and *bekanntlich* ('as is well known'). Intuitively, these markers differ in how they allow a speaker to present p as settled in a given common ground. We ask whether this intuition can be formalized in a probabilistic pragmatic model that infers marker-specific felicity thresholds from production choices and uses them to predict marker choice and listener interpretation as controversy, community, and goal vary.

- (1) *Impfungen sind (soviel ich weiß/ ja / bekanntlich) sicher und wirksam; deshalb solltest du dich impfen lassen.*

'Vaccinations are, as far as I know / as you know / as is well known, safe and effective; therefore you should get vaccinated.'

With *soviel ich weiß*, the speaker restricts p to their own epistemic state; *ja* presents p as shared with the addressee; and *bekanntlich* presents p as more broadly publicly accessible knowledge. The felicity of each marker therefore depends on several dimensions at once. First, the status of p is community-relative: the safety and efficacy of vaccination may be treated as settled in a medical-expert community, but remain contested in a divided public setting. Second, the speaker's goal matters. A political group promoting vaccination, for instance, may have an interest in presenting p as already settled, because doing so strengthens the move from p to the recommendation q . Third, propositions differ in their baseline degree of controversy, independently of the immediate discourse context. Thus, by selecting a particular consensus marker, the speaker can present p as more or less settled and thereby guide the addressee toward accepting q as a recommendation grounded in shared assumptions. We use the label *consensus markers* functionally: the three expressions differ in syntactic category and lexical meaning, but all bear on how p is managed in the common ground. Previous work on German discourse particles treats *ja* as contributing non-truth-conditional, use-conditional meaning [1, 2]. Its common-ground effect has been characterized as marking p as KNOWN [3], speaker-assumed common knowledge [4], uncontroversial [5–7], settled in the common ground [8], or available for manipulative use in convincing an interlocutor [9]. Recent experimental work further links *ja* to causal discourse contexts [10], motivating the running setup in (1), where p serves as the premise for the downstream recommendation q . What remains missing is a formal account of marker felicity that integrates community-relative consensus around p with the speaker's communicative goal in presenting p as support for q .

Building on these insights, we propose a utility-based, noisy-threshold model within probabilistic pragmatics [11–13]. The central modelling contribution is to infer these felicity thresholds from production data and use posterior predictions for item selection of listener's tasks to connect production and interpretation. Marker choice is modelled as a decision about whether p can be presented as a sufficiently settled premise for q . This brings together the propositional controversy of p , the pragmatic status of p in a relevant epistemic community, and the speaker's goal in using p to support q . Empirically, we implement this link in a three-step pipeline (Fig. 2): norming estimates baseline propositional controversy, production estimates how pragmatic controversy and goal strength shape marker choice, and comprehension tests whether listeners recover speaker goals and adoption consequences from the chosen marker. The empirical findings are graded in both production and interpretation: lower controversy and higher goal strength favour stricter thresholds, and the corresponding markers lead listeners to infer stronger speaker goals. For adoption, weakly supported felicity conditions can make the speaker's attempt to present p as settled visible, leading listeners to discount the recommendation, in line with pragmatic reactance [14]. Fitted to the production data, the model estimates partially overlapping felicity thresholds for *ja* and *bekanntlich* from speakers' choices across contexts. Taken together, the results show that consensus markers constrain how p can be presented as a shared premise for q , with speaker choice and listener inference shaped by the relevant epistemic community and the speaker's goal.

Model.¹ Perceived controversy combines propositional and pragmatic components, $pc = pc_{\text{prop}} \cdot pc_{\text{prag}}$, and enters a utility $U(pc, g) = -w_{pc} \cdot pc + w_g \cdot g - w_{\text{int}} \cdot pc \cdot g$, where g is the speaker’s goal strength. Low controversy and high goal strength jointly favour stricter thresholds, while w_{int} captures the additional cost of invoking consensus under high controversy. Each marker u_i is associated with a noisy interval $[\theta_i, \theta_{i+1})$ on the utility scale, so relative position and overlap are inferred from choices. A marker is selected when U falls within its interval, $P(u_i | pc, g) \propto \sigma(\lambda(U - \theta_i)) \cdot \sigma(-\lambda(U - \theta_{i+1}))$, and a pragmatic listener inverts this model to infer pc and g .

Empirical evidence.² The experiments implement the pipeline summarized above (Fig. 2). In production (Fig. 1c), g favours stricter-threshold markers (0.28, $P(> 0) = 1$) and pc_{prag} disfavors them (-0.17 , $P(< 0) = 0.98$); the goal effect shrinks under high controversy (interaction: -0.22 , $P(< 0) = 0.97$), as predicted by the w_{int} cost term. In comprehension (Fig. 1d), both ja (0.49, $P(> 0) = 1$) and $bekanntlich$ (0.46, $P(> 0) = 1$) increase inferred g ; for adoption, ja shows a credible positive effect (0.42, $P(> 0) = 0.99$), and $bekanntlich$ shows a weaker one (0.17, $P(> 0) = 0.90$). This dissociation is compatible with pragmatic reactance [14]: listeners infer the speaker’s goal from marker choice, then discount weakly supported consensus marking as argumentative pressure. Our production-trained noisy-threshold model captures these patterns within a single framework.




Figure 1: Model and experimental results. Points show means; error bars show bootstrap 95% confidence intervals.

¹Implemented in Stan; MCMC inference used the No-U-Turn Sampler [15, 16].

²Reported coefficients and one-sided probabilities come from Bayesian regression models; hypotheses were evaluated using posterior predictions.




Figure 2: Experimental pipeline. (A) Participants discuss social topics with a partner, with p serving as support for a recommendation q . (B) Sequential design: a norming study establishes baseline propositional controversy, followed by production and comprehension experiments. (C) Norming: eight topics rated for propositional controversy. (D) Production: participants select the best consensus marker given manipulated pragmatic controversy and goal strength. (E) Comprehension: participants infer speaker goal strength and adoption likelihood from marker-containing utterances.

References:

- [1] Daniel Gutzmann. *Use-conditional meaning: Studies in multidimensional semantics*, volume 6. OUP Oxford, 2015.
- [2] Patrick Grosz. Discourse particles. *The Wiley Blackwell companion to semantics*, pages 619–652, 2021.
- [3] Maria Thurmair. *Modalpartikeln und ihre Kombinationen*. Niemeyer, 1989.
- [4] Elena Karagjosova. *The meaning and function of German modal particles*. DFKI, 2004.
- [5] Joachim Jacobs. On the semantics of modal particles. In Werner Abraham, editor, *Descriptive and theoretical investigations on the logical, syntactic and pragmatic properties of discourse particles in German*, pages 141–162. John Benjamins Publishing Company, 1991. doi: 10.1075/pbns.12.06jac.
- [6] Malte Zimmermann. Discourse particles. In Claudia Maienborn, Klaus von Heusinger, and Paul Portner, editors, *Semantics: An International Handbook of Natural Language Meaning*. Mouton de Gruyter, Berlin, 2012.
- [7] Richard Walerit. Modal particles and their functional equivalents: A speech-act-theoretic approach. *Journal of Pragmatics*, 33(9):1391–1417, 2001.
- [8] Sophie Repp. Common ground management: Modal particles, illocutionary negation and verum. In *Beyond expressives: Explorations in use-conditional meaning*, pages 231–274. Brill, 2013.
- [9] Sophia Döring. *Modal Particles, Discourse Structure and Common Ground Management: Theoretical and Empirical Aspects*. PhD thesis, Sprach- und literaturwissenschaftlichen Fakultät, Humboldt-Universität zu Berlin, Berlin, 2016.
- [10] Hannah Seemann. *The Signaling Function of Modal Particles in Discourse*. PhD thesis, Ruhr-Universität Bochum, 2025. URL <https://osf.io/tw3mb>.
- [11] Noah D Goodman and Michael C Frank. Pragmatic language interpretation as probabilistic inference. *Trends in cognitive sciences*, 20(11):818–829, 2016.
- [12] Michael Franke and Gerhard Jäger. Probabilistic pragmatics, or why Bayes’ rule is probably important for pragmatics. *Zeitschrift für Sprachwissenschaft*, 35(1): 3–44, 2016.
- [13] Judith Degen. The rational speech act framework. *Annual Review of Linguistics*, 9(1):519–540, 2023.
- [14] Jack Williams Brehm. *A Theory of Psychological Reactance*. Academic Press, New York, 1966.
- [15] Bob Carpenter, Andrew Gelman, Matthew D. Hoffman, Daniel Lee, Ben Goodrich, Michael Betancourt, Marcus Brubaker, Jiqiang Guo, Peter Li, and Allen Riddell. Stan: A probabilistic programming language. *Journal of Statistical Software*, 76(1):1–32, 2017. doi: 10.18637/jss.v076.i01.
- [16] Matthew D. Hoffman and Andrew Gelman. The No-U-Turn sampler: Adaptively setting path lengths in Hamiltonian Monte Carlo. *Journal of Machine Learning Research*, 15(47):1593–1623, 2014. URL <https://jmlr.org/papers/v15/hoffman14a.html>.